

## Module - 4

### Transport Layer:

The Transport Layer is one of the seven layers of the OSI (Open Systems Interconnection) model, and is responsible for end-to-end communication and data transfer between devices. It ensures reliable and orderly data delivery by breaking data into smaller segments and reassembling them at the destination. The main protocols in this layer are TCP (Transmission Control Protocol) and UDP (User Datagram Protocol).

1. **Segmentation and Reassembly:** The Transport Layer divides data into smaller units called segments. This segmentation is necessary because data generated by applications can be larger than what can be transmitted efficiently over a network. The Transport Layer on the receiving end reassembles these segments into the original data stream.
2. **End-to-End Communication:** The Transport Layer provides end-to-end communication, meaning it ensures that data sent from one device reaches its intended destination device accurately and in the correct order. It abstracts the complexities of the underlying network and shields the higher-layer applications from network-specific details.
3. **Flow Control:** Flow control mechanisms in the Transport Layer prevent congestion and ensure that data is transmitted at a rate that the receiving device can handle. This prevents overwhelming the receiver with data, leading to packet loss or network congestion.
4. **Error Detection and Correction:** The Transport Layer may implement error detection and correction mechanisms to ensure data integrity during transmission. This is particularly important for reliable communication.
5. **Multiplexing and Demultiplexing:** Multiplexing allows multiple applications or sessions to share the same network connection or port on a device. Demultiplexing, on the receiving end, ensures that incoming data is correctly routed to the appropriate application or session.
6. **Port Numbers:** Port numbers are used to identify specific services or applications on a device. Together with IP addresses, port numbers enable devices to determine which application should receive incoming data. For example, web traffic typically uses port 80 for HTTP and port 443 for HTTPS.
7. **Transport Layer Protocols:** The Transport Layer supports multiple protocols, with the two most common ones being TCP (Transmission Control Protocol) and UDP (User Datagram Protocol). TCP provides reliable, connection-oriented communication with error recovery and flow control, making it suitable for applications like web browsing and file transfers. UDP, on the other hand, offers a connectionless, lightweight, and low-latency communication method, making it suitable for applications like real-time audio and video streaming.
8. **Session Establishment and Termination:** In the case of TCP, the Transport Layer handles the establishment, maintenance, and termination of sessions (connections) between devices. This ensures that data is reliably delivered and that both sender and receiver are in sync during communication.

### **Transport layer services provided to the upper layers elements of transport protocol addressing connection establishment, Connection release, Error Control & Flow Control.**

The Transport Layer provides critical services to the upper layers (Layers 5 to 7 in the OSI model) of the networking stack, ensuring that data is efficiently and reliably transmitted between applications or processes running on different devices in a network. Two fundamental services provided by the Transport Layer are transport protocol addressing and connection establishment:

## 1. Transport Protocol Addressing:

- Port Numbers: Port numbers are a fundamental aspect of transport protocol addressing. They are used to distinguish different services or applications running on a single device. Port numbers are 16-bit values, allowing for up to 65,535 unique port assignments. These assignments are standardized for common services (e.g., HTTP uses port 80, HTTPS uses port 443).

- Source and Destination Ports: In the Transport Layer header, each packet or segment includes source and destination port numbers. These port numbers are crucial for demultiplexing incoming data on the receiving end. They ensure that data is correctly routed to the appropriate application or process based on the port numbers.

## 2. Connection Establishment:

- Connection-Oriented Protocols: Some transport layer protocols, such as TCP (Transmission Control Protocol), provide connection-oriented communication. Connection establishment involves a series of steps:

- SYN (Synchronize) Phase: The sender (client) initiates a connection by sending a SYN packet to the receiver (server).

- SYN-ACK Phase: The receiver acknowledges the request by sending a SYN-ACK packet back to the sender.

- ACK (Acknowledgment) Phase: The sender acknowledges the acknowledgment by sending an ACK packet to the receiver. At this point, the connection is established.

- Reliability and Error Handling: Connection-oriented protocols like TCP ensure reliable data delivery. They establish a connection, provide mechanisms for error detection and correction, and manage the flow of data to prevent congestion and data loss.

These services provided by the Transport Layer play a crucial role in ensuring that data is transmitted accurately and efficiently between applications or processes running on different devices:

- End-to-End Communication: The use of port numbers and connection establishment ensures that data is delivered accurately and reliably to the intended application or process on the receiving device, regardless of the complexities of the network in between.

- Multiplexing: Port numbers allow multiple applications or sessions to share the same network connection or network interface on a device. This efficient use of network resources is essential for modern networking, where many services coexist on a single device.

- Flow Control: Flow control mechanisms help prevent network congestion and ensure that data is transmitted at a rate that the receiving device can handle. This prevents data loss and maintains the quality of service.

- Error Handling and Correction: Transport layer protocols often include error detection and correction mechanisms to ensure the integrity of data during transmission.

## Internet Transport Protocols:

### UDP vs TCP:

#### TCP:

TCP, or Transmission Control Protocol, is one of the core protocols of the Internet Protocol suite and operates at the transport layer (Layer 4) of the OSI model. It provides reliable, connection-oriented communication between devices on a network. Here are key aspects and features of TCP:

1. Connection-Oriented: TCP is a connection-oriented protocol, which means it establishes a logical connection between the sender and receiver before data exchange begins. This connection is established using a process called the "three-way handshake."

2. Reliability: TCP ensures reliable data delivery. It achieves this by using a combination of mechanisms:

- Acknowledgment (ACK): The receiver acknowledges the receipt of each data segment.
- Retransmission: If a segment is not acknowledged within a certain time, TCP retransmits it.
- Sequence Numbers: Each data segment is assigned a sequence number, allowing the receiver to reconstruct the data in the correct order.
- Checksums: TCP uses checksums to detect errors in the data, and if errors are detected, it requests retransmission.

3. Flow Control: TCP incorporates flow control mechanisms to prevent congestion and manage the rate of data transfer. It ensures that data is sent at a pace the receiving device can handle, preventing packet loss and network congestion.

4. Ordered Delivery: TCP guarantees that data is delivered to the application in the same order it was sent. This is crucial for applications that rely on the correct sequence of data, such as web pages, file transfers, and email.

5. Full Duplex Communication: TCP supports full-duplex communication, allowing data to be transmitted in both directions simultaneously. This is achieved through separate send and receive buffers.

6. High Overhead: TCP headers contain additional information, including sequence numbers, acknowledgment numbers, and control flags. This results in a relatively higher overhead compared to some other transport layer protocols like UDP (User Datagram Protocol).

7. Connection Termination: TCP handles the graceful termination of a connection through a process known as the "four-way handshake," ensuring that any remaining data is transmitted and received before the connection is closed.

8. Applications: TCP is commonly used for applications that require reliable and ordered data delivery, such as web browsing, email, file transfers, remote desktop connections, and any scenario where data integrity is critical.

9. Port Numbers: TCP uses port numbers to identify specific services or applications on devices. Port numbers, combined with IP addresses, enable devices to determine which application should receive incoming data.

## **UDP:**

UDP, or User Datagram Protocol, is a core protocol of the Internet Protocol suite and operates at the transport layer (Layer 4) of the OSI model. Unlike TCP (Transmission Control Protocol), UDP is a connectionless and lightweight transport protocol that provides minimal services. Here are key aspects and features of UDP:

1. Connectionless: UDP is a connectionless protocol, which means it does not establish a dedicated connection before sending data. It simply sends datagrams (packets) independently to the recipient without prior setup.

2. Minimal Overhead: UDP has minimal header overhead compared to TCP. This makes it a lightweight protocol and results in lower latency for data transmission.

3. Unreliable: Unlike TCP, UDP does not provide reliability features such as acknowledgments, retransmissions, and error correction. Once a UDP packet is sent, the sender does not know whether it was successfully received by the recipient.

4. No Flow Control: UDP does not incorporate flow control mechanisms to manage the rate of data transfer or prevent congestion. This means that UDP can transmit data at its maximum rate, potentially leading to network congestion if not managed properly.

5. No Ordered Delivery: UDP makes no guarantee about the order in which packets are delivered. Packets can arrive at the recipient out of order compared to how they were sent.

6. Broadcast and Multicast Support: UDP can be used for broadcasting data to multiple recipients simultaneously (broadcast) or sending data to specific groups of recipients (multicast). This makes UDP suitable for real-time multimedia streaming and online gaming.

7. Low Latency: Due to its minimal overhead and lack of handshakes, UDP offers lower latency compared to TCP. This makes it ideal for real-time applications where low delay is critical, such as voice and video conferencing.

8. Examples of Use: UDP is commonly used in scenarios where speed and low overhead are more important than reliability and error correction. Examples include real-time multimedia streaming (e.g., VoIP and video conferencing), online gaming, DNS (Domain Name System) queries, and DHCP (Dynamic Host Configuration Protocol) for IP address assignment.

9. Port Numbers: Like TCP, UDP also uses port numbers to identify specific services or applications on devices. Port numbers, combined with IP addresses, enable devices to determine which application should receive incoming data.

### **The TCP Service Model:**

The TCP (Transmission Control Protocol) service model defines a set of features and characteristics that TCP provides to applications and users when transmitting data over a network. TCP is a connection-oriented and reliable transport protocol, and its service model reflects these key attributes:

1. Connection-Oriented Communication: TCP establishes a logical connection between the sender and receiver before data exchange begins. This connection is established through a process called the "three-way handshake." The connection ensures that both parties are synchronized and ready for data transmission.

2. Reliable Data Delivery: TCP guarantees reliable data delivery. It achieves this reliability through various mechanisms:

- Acknowledgment (ACK): TCP uses acknowledgments to confirm the successful receipt of data segments. When the sender transmits data, the receiver sends ACKs back to indicate that the data arrived without errors.
- Retransmission: If an ACK is not received within a certain time or if data is lost or corrupted during transmission, TCP retransmits the data segments until they are successfully received and acknowledged.
- Sequence Numbers: TCP assigns a unique sequence number to each data segment, allowing the receiver to reconstruct the data in the correct order.
- Checksums: TCP uses checksums to detect errors in the data. If errors are detected, TCP requests retransmission.

3. Flow Control: TCP incorporates flow control mechanisms to manage the rate of data transfer and prevent network congestion. Flow control ensures that data is sent at a pace that the receiving device can handle. It prevents situations where the sender overwhelms the receiver with data, leading to packet loss or network congestion.

4. Ordered Delivery: TCP guarantees that data is delivered to the application in the same order it was sent. This is essential for applications that rely on the correct sequence of data, such as web pages, file transfers, and email.

5. Full Duplex Communication: TCP supports full-duplex communication, allowing data to be transmitted in both directions simultaneously. This is achieved through separate send and receive buffers, enabling bi-directional data exchange without interference.

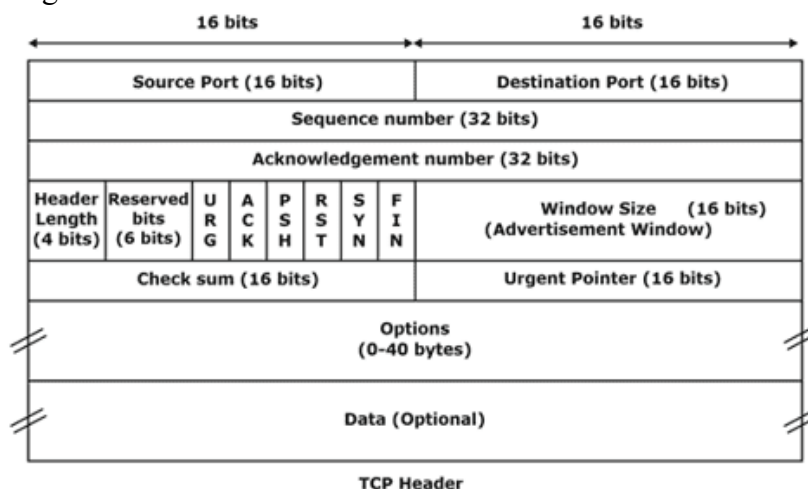
6. Error Detection and Correction: TCP includes error detection mechanisms to ensure data integrity during transmission. If errors are detected, TCP can request retransmission of the affected data.

7. Port Numbers: TCP uses port numbers to identify specific services or applications on devices. Port numbers, combined with IP addresses, enable devices to determine which application should receive incoming data.

8. Connection Termination: TCP handles the graceful termination of a connection through a process known as the "four-way handshake." This ensures that any remaining data is transmitted and received before the connection is closed.

## The TCP Segment Header:

The TCP (Transmission Control Protocol) segment header is a crucial part of the TCP protocol, used for packaging and transmitting data over a network. It contains various fields and information necessary for the proper delivery and management of data between the sender and receiver. Here's an explanation of the key fields in the TCP segment header:



1. Source Port (16 bits): This field specifies the port number of the sender's application or process. It identifies the source of the TCP segment on the sender's side.

2. Destination Port (16 bits): This field specifies the port number of the receiver's application or process. It identifies the intended destination for the TCP segment on the receiver's side.

3. Sequence Number (32 bits): The sequence number field is used for data ordering and reassembly. It assigns a unique sequence number to each byte of data sent. This enables the receiver to reconstruct the data in the correct order.

4. Acknowledgment Number (32 bits): This field is used to acknowledge the receipt of data segments. It contains the next expected sequence number the receiver is anticipating. When a receiver acknowledges a segment, it indicates that all data up to that sequence number has been successfully received.

5. Data Offset (4 bits): The data offset field specifies the length of the TCP header in 32-bit words. It indicates where the data begins within the TCP segment.

6. Reserved (6 bits): These bits are reserved for future use and should be set to zero.

7. Flags (6 bits): The flags field contains various control flags that determine the purpose and behavior of the TCP segment. Some important flags include:

- URG (Urgent Pointer): Indicates urgent data in the segment.
- ACK (Acknowledgment): Indicates that the acknowledgment field is significant.
- PSH (Push): Indicates that the receiver should push data to the application as soon as possible.
- RST (Reset): Resets the connection.

- SYN (Synchronize): Initiates a connection establishment.
- FIN (Finish): Initiates a connection termination.

8. Window Size (16 bits): The window size field indicates the size of the sender's receive window. It represents the amount of data, in bytes, that the sender is willing to receive before requiring acknowledgment from the receiver.

9. Checksum (16 bits): The checksum field contains a checksum value computed over the TCP header, data, and a pseudo-header (including source and destination IP addresses). It is used to detect errors in the TCP segment during transmission.

10. Urgent Pointer (16 bits): If the URG flag is set, the urgent pointer field specifies an offset from the sequence number indicating the last urgent data byte in the segment. It is used to indicate urgent data that requires immediate attention.

11. Options (variable length): The options field is optional and can contain various TCP options, such as Maximum Segment Size (MSS), Timestamps, and Window Scale. These options provide additional control and tuning capabilities for TCP.

12. Padding (variable length): Padding is added to align the TCP header to a 32-bit boundary when necessary. It ensures that the header length is a multiple of 32 bits.

The TCP segment header, along with the TCP service model and reliable data delivery mechanisms, ensures that data is transmitted accurately and efficiently between applications or processes across a network. It plays a vital role in managing connections, sequencing data, and detecting and recovering from errors during transmission.

## **Connection Establishment:**

Connection establishment in networking, especially in the context of the TCP (Transmission Control Protocol) protocol, is the process by which two devices or endpoints establish a communication link or connection before they can exchange data. This process ensures that both parties are synchronized and ready to transmit and receive data reliably. The connection establishment process typically involves a series of steps, often referred to as the "three-way handshake."

Here's a high-level overview of the connection establishment process in TCP:

### **1. Step 1: Initiating the Connection (SYN):**

- The process begins when one device, often referred to as the client, wants to establish a connection with another device, known as the server.
- The client sends a TCP segment with the SYN (Synchronize) flag set. This initial segment is sometimes called the "SYN segment."
- The SYN segment contains the client's initial sequence number (ISN), which is a random value chosen by the client to identify the starting point for data sequencing.

### **2. Step 2: Acknowledging the Connection Request (SYN-ACK):**

- Upon receiving the SYN segment, the server acknowledges the connection request by sending a TCP segment with both the SYN and ACK flags set. This is called the "SYN-ACK segment."
- The SYN-ACK segment also contains the server's initial sequence number (ISN) and acknowledges the client's ISN by incrementing it by one.

### **3. Step 3: Finalizing the Connection (ACK):**

- The client receives the SYN-ACK segment from the server and acknowledges it by sending another TCP segment with the ACK flag set. This segment is sometimes called the "ACK segment."
- The ACK segment acknowledges the server's ISN by incrementing it by one.

At this point, the TCP connection is considered established, and both the client and server can begin transmitting data reliably in both directions. The connection establishment process ensures that both parties are aware of each other's readiness to communicate and that initial sequence numbers are synchronized for data sequencing and acknowledgment purposes.

The three-way handshake in TCP is designed to be a reliable and robust method for establishing connections while ensuring that both parties are in agreement about the connection's parameters. It helps prevent issues like data misordering and protects against unauthorized or malicious connection attempts.

When the communication session is complete, a similar process, often called the "four-way handshake," is used to gracefully terminate the connection and ensure that any remaining data is transmitted and received before the connection is closed.

## **Connection Release:**

Connection release, also known as connection termination, is the process by which two devices or endpoints terminate an established communication link or connection in a network. In the context of the TCP (Transmission Control Protocol) protocol, connection release is often referred to as the "four-way handshake" because it typically involves a series of four steps to ensure that both parties agree to end the connection gracefully.

Here's an overview of the connection release process in TCP:

### **1. Step 1: Initiating the Connection Release (FIN):**

- One of the devices, either the client or the server, initiates the connection release by sending a TCP segment with the FIN (Finish) flag set. This segment is called the "FIN segment."
- The FIN segment indicates that the sender has finished sending data and wants to close the connection. However, it may still be willing to receive data.

### **2. Step 2: Acknowledging the Release Request (ACK):**

- Upon receiving the FIN segment, the receiving device acknowledges the release request by sending an ACK (Acknowledgment) segment back to the sender. This ACK segment acknowledges the FIN segment.
- At this point, the receiving device is still allowed to send data, but it agrees to stop sending data once it has none left to transmit.

### **3. Step 3: Releasing Data (FIN):**

- If the receiving device has any remaining data to send, it can continue to do so until it is ready to release the connection.
- When the receiving device is ready to release the connection, it sends its own FIN segment, indicating that it has finished sending data. This is another FIN segment but from the other party.

### **4. Step 4: Final Acknowledgment (ACK):**

- Upon receiving the second FIN segment, the sender acknowledges it with an ACK segment. This ACK segment acknowledges the second FIN segment.
- Once this final ACK is sent and acknowledged, both parties consider the connection released, and the connection is closed.

The four-way handshake ensures that both the sender and receiver have agreed to close the connection and that any remaining data is transmitted and received before termination. This process helps prevent issues like data loss due to abrupt connection termination and ensures a graceful release of the connection.

Once the connection is closed, neither party can transmit data on that connection, and the resources associated with the connection are freed up for use elsewhere. Connection release is an essential part of reliable and orderly data transmission over TCP connections.

## **TCP Sliding Window:**

TCP (Transmission Control Protocol) sliding window is a fundamental mechanism that plays a crucial role in optimizing the efficiency and reliability of data transmission over a TCP connection. It is used to manage the flow of data between the sender and receiver while taking into account network conditions, receiver's capabilities, and congestion control. Here's an explanation of TCP sliding window:

### **1. Overview of Sliding Window:**

- **Window Size:** In the context of TCP, a "window" is a range of sequence numbers that the sender is allowed to transmit without receiving an acknowledgment from the receiver. This window size is dynamic and can vary during the course of the connection.

### **2. Sender's Perspective:**

- **Sender's Window:** The sender maintains a sliding window of data it is allowed to send. The size of this window, known as the "send window" or "sender's window," is determined by several factors, including the receiver's advertised window size and network conditions.

- **Sequence Numbers:** Each segment sent by the sender is assigned a unique sequence number. The sender's window represents the range of acceptable sequence numbers that the sender can transmit. The sender sends data within this range.

- **Acknowledgments:** The sender keeps track of which segments have been acknowledged by the receiver. As acknowledgments arrive, the sender adjusts the window to allow for the transmission of new data.

- **Flow Control:** The sender uses the sliding window to implement flow control, ensuring it does not overwhelm the receiver with data. It only sends data that falls within the current window.

### **3. Receiver's Perspective:**

- **Receiver's Window:** The receiver maintains a corresponding window on its side, known as the "receive window" or "receiver's window." This window indicates the range of sequence numbers it is willing to accept from the sender.

- **Advertised Window:** The receiver periodically informs the sender of its receive window size through TCP acknowledgment (ACK) segments. This is known as the "advertised window." The sender uses the advertised window size to determine how much data it can send before waiting for acknowledgments.

### **4. Dynamic Adjustment:**

- **Sliding Mechanism:** The sliding window mechanism allows the sender and receiver to dynamically adjust their windows based on several factors, including network conditions and the receiver's ability to process data.

- **Efficient Use of Bandwidth:** TCP sliding window ensures that the sender keeps the network busy with a reasonable amount of unacknowledged data, maximizing network utilization while preventing congestion.

### **5. Congestion Control:**

- **TCP sliding window** is closely related to congestion control mechanisms. If the sender observes congestion (e.g., due to packet loss), it may reduce the size of its window to alleviate congestion and retransmit lost segments.



## 6. Buffer Management:

- Sliding window also helps manage the sender's and receiver's buffer space efficiently. The sender can transmit data that fits within the receiver's window, and the receiver can buffer incoming data based on its advertised window size.

In summary, TCP sliding window is a dynamic mechanism that allows the sender and receiver to efficiently manage the flow of data in a TCP connection. It ensures that data is transmitted at a rate that the network and receiver can handle while maximizing throughput and reliability. The sliding window's size and position adapt in real-time to network conditions, providing optimal performance for data transmission.

## Congestion Control Algorithm:

Congestion control algorithms in computer networks are mechanisms and strategies used to manage and alleviate network congestion. Network congestion occurs when the demand for network resources, such as bandwidth, exceeds the available capacity, leading to performance degradation, packet loss, and increased latency. Congestion control algorithms aim to prevent or mitigate these issues by regulating the flow of data through the network. Here's an explanation of congestion control algorithms:

### 1. Detection of Congestion:

- Congestion control algorithms start by detecting signs of congestion in the network. Common indicators include increased packet loss, high network utilization, and rising packet queuing delays at routers.

### 2. Feedback Mechanisms:

- Many congestion control algorithms rely on feedback mechanisms to gather information about network conditions. This feedback can come from routers, switches, or other network devices. Some algorithms also use end-to-end feedback from receivers or other endpoints.

### 3. Congestion Avoidance:

- Congestion control algorithms aim to avoid congestion by adjusting the rate at which data is sent into the network. This is often done through rate or window size adjustments.

### 4. Slow Start:

- Some algorithms, like TCP's congestion control, use a "slow start" mechanism when establishing a connection. During this phase, the sender gradually increases the rate at which it sends data until congestion is detected.

### 5. Congestion Detection and Reaction:

- When congestion is detected or inferred from feedback, congestion control algorithms take action to reduce the data transmission rate. This can involve reducing the sending rate, dropping or marking packets, or applying other congestion mitigation strategies.

### 6. Adaptive Rate Control:

- Many modern congestion control algorithms are adaptive, meaning they continuously monitor network conditions and adjust their sending rates dynamically. They aim to find an optimal sending rate that maximizes network utilization without causing congestion.

### 7. Fairness:

- Congestion control algorithms often strive to be fair, ensuring that multiple flows or users share the available network resources fairly. Fairness mechanisms prevent a single flow from dominating the network and degrading the performance of others.

### 8. Explicit Congestion Signaling:

- Some algorithms use explicit signals, like Explicit Congestion Notification (ECN), to inform endpoints of congestion. Routers may mark packets or send ECN signals to indicate congestion, and endpoints adjust their behavior accordingly.

9. Random Early Detection (RED):

- RED is a router-based congestion control mechanism that proactively drops or marks packets before congestion becomes severe. It aims to maintain a moderate level of queuing delay and avoid the "congestion collapse" phenomenon.

10. Window-Based Control:

- Some algorithms, like TCP congestion control, use a window-based approach to control congestion. The sender maintains a congestion window that limits the number of unacknowledged packets in the network.

11. Explicit Rate Control:

- Some congestion control algorithms explicitly regulate the sending rate based on network conditions. These algorithms aim to determine the "right" sending rate for optimal network performance.

Effective congestion control is essential for ensuring that networks operate efficiently, deliver low-latency communication, and provide fair access to resources for all users and applications. Different congestion control algorithms may be suitable for different network environments and use cases, and ongoing research continues to improve these algorithms' performance and adaptability.

## UDP Header Format

