

## Module-5

**Views:** A view is a virtual or saved query result that behaves like a table. A view allows you to encapsulate complex SQL queries into a reusable object, making it easier to retrieve and work with specific sets of data.

Views are a powerful tool in SQL that enhance data abstraction, simplify query complexity, and improve security and maintainability in database systems. They are especially useful when dealing with complex database schemas or when you want to provide controlled access to specific subsets of data.

A "view" is a virtual table that is created by executing a SELECT query. Unlike physical tables, views do not store data themselves; they are dynamic, stored queries that generate a result set whenever they are queried. Views are used to simplify data access, enhance security, and provide a convenient way to work with complex or frequently used queries.

### **Views provide the following benefits:**

1. **Abstraction:** Views abstract the underlying data model by allowing users to interact with a simplified, virtual representation of the data. Users can query a view without needing to understand the complex SQL logic or structure of the underlying tables.
2. **Security:** Views can be used to enforce security and access control by restricting the columns and rows that users can access. This can help protect sensitive data by only exposing the necessary information to authorized users.
3. **Simplicity:** Views can simplify complex joins and calculations. Instead of writing complex queries repeatedly, you can create a view with the necessary logic, and then query the view as if it were a regular table.
4. **Reusability:** Views promote code reusability. Once a view is created, it can be used in multiple queries, reports, or applications, reducing the need to rewrite the same logic.

Here's how you create and use views in SQL:

### **Creating a View:**

To create a view, you use the `CREATE VIEW` statement followed by the view name and the SQL query that defines the view. For example:

```
CREATE VIEW employee_view AS  
SELECT employee_id, first_name, last_name, department  
FROM employees  
WHERE department = 'HR';
```

### **Querying a View:**

Once a view is created, you can query it just like you would a regular table:

```
SELECT FROM employee_view;
```

This query retrieves data from the `employee_view` view.

### **Modifying a View:**

You can also modify views. To do this, you can use the `ALTER VIEW` statement to change the underlying SQL query of the view. Be cautious when modifying views, as any changes may affect queries that rely on the view.

### **Dropping a View:**

To remove a view from the database, you can use the `DROP VIEW` statement:

```
DROP VIEW employee_view;
```

### **Create view:**

To create a view in SQL, you use the `CREATE VIEW` statement followed by the view name and the SQL query that defines the view. Here's the general syntax for creating a view:

```
CREATE VIEW view_name AS  
SELECT column1, column2, ...  
FROM table_name  
WHERE conditions;
```

You can create a view by specifying the columns you want to include in the view and defining the query that retrieves the data. Here's an example of how to create a view:

Suppose you have a table called `employees` with the following structure:

```
CREATE TABLE employees (  
  employee_id INT PRIMARY KEY,  
  first_name VARCHAR(50),  
  last_name VARCHAR(50),  
  department VARCHAR(50),  
  salary DECIMAL(10, 2)  
);
```

You can create a view that displays the employee IDs and full names of employees in the HR department:

```
CREATE VIEW hr_employees AS  
SELECT employee_id, CONCAT(first_name, ' ', last_name) AS full_name  
FROM employees  
WHERE department = 'HR';
```

In this example, the `hr\_employees` view is created with two columns: `employee\_id` and `full\_name`. It retrieves data from the `employees` table, concatenating the `first\_name` and `last\_name` columns to create the `full\_name` column, and filters the results to include only employees in the 'HR' department.

You can then query this view like you would a regular table:

```
SELECT FROM hr_employees;
```

This query will retrieve the employee IDs and full names of HR department employees from the `hr\_employees` view.

Keep in mind that views are virtual tables that don't store data themselves; they reflect the data in the

underlying tables. Therefore, any changes made to the underlying tables will be reflected in the view's results.

To drop (delete) a view, you can use the `DROP VIEW` statement:

**DROP VIEW hr\_employees;**

This will remove the `hr\_employees` view from the database.

### Types of view:

There are two main types of views in SQL:

- **Materialized views:** Materialized views are physical tables that contain the data that is specified in the CREATE VIEW statement. Materialized views are updated automatically when the underlying tables are updated.
- Materialized views are useful for improving the performance of queries that frequently access a large amount of data. This is because materialized views are stored in the database and can be accessed directly, without having to query the underlying tables.
- **Non-materialized views:** Non-materialized views are virtual tables that do not contain any data. Non-materialized views are created when they are queried and are updated automatically when the underlying tables are updated.

### Operations on views:

Views in SQL are virtual tables that provide a way to access and manipulate data from one or more underlying tables. While views are often used for querying and simplifying data access, they can also be used for various operations such as updating, inserting, and deleting data under certain conditions. Here are the operations you can perform on views:

#### 1. SELECT Operation:

- The primary operation on a view is selecting data. You can query a view just like you would a regular table using the `SELECT` statement. Views allow you to present data from one or more tables in a simplified and customized manner.

**SELECT FROM my\_view;**

#### 2. INSERT Operation:

- Some views, known as updatable views, allow you to insert data into the underlying tables through the view. To perform an `INSERT` operation on a view, the view must meet specific criteria:

- The view must be based on a single table.
- The view's columns must map directly to the columns in the underlying table.
- The view must not include any computed columns or expressions.

**INSERT INTO my\_view (column1, column2) VALUES (value1, value2);**

#### 3. UPDATE Operation:

- Updatable views also allow you to update data in the underlying tables. Like with insert operations, there are specific conditions that must be met for an update operation to work on a view.

**UPDATE my\_view SET column1 = new\_value WHERE condition;**

#### 4..Dropping Views:

- You can use the 'DROP VIEW' statement to delete a view when it is no longer needed.

**DROP VIEW my\_view;**

#### 5. DELETE Operation:

- Similarly, updatable views can be used to delete data from the underlying tables.

**DELETE FROM my\_view WHERE condition;**

#### 6. Joins and Aggregations:

- Views can be used as part of more complex queries, including joins with other tables and aggregations. This allows you to work with views alongside other tables to retrieve and analyze data.

**SELECT FROM my\_view  
JOIN another\_table ON my\_view.column = another\_table.column;**

#### Dictionaries for table:

In the context of databases and SQL, dictionaries or data dictionaries are not typically associated with tables. However, you might be referring to something related to tables or database schema documentation. Let me clarify a few concepts:

##### 1. Data Dictionary:

- A data dictionary is a repository of metadata or information about the data within a database. It describes the structure of the database, including tables, columns, data types, constraints, and relationships.
- In SQL databases, data dictionaries are often system tables or views that store information about the database schema. For example, in MySQL, you can query the 'information\_schema' database to retrieve information about tables and columns.

##### 2. Table:

- In a database, a table is a structured collection of data organized into rows and columns. Each column represents an attribute or field, and each row represents a record or entry.
- Tables are used to store, retrieve, and manipulate data in a database. They have names and are defined by their schema, which includes column names, data types, and constraints.

##### 3. Dictionary Tables (or System Tables):

- Some database systems include special system tables or views that provide metadata about the database objects, including tables. These system tables can be queried to obtain information about the database schema.
- For example, in Oracle, you can query the 'ALL\_TABLES' or 'USER\_TABLES' views to retrieve information about user tables in the database.

If you are looking to document the structure of your database tables, including their columns, data types, constraints, and relationships, you can create your own data dictionary or use database documentation tools to generate documentation from the database's metadata. These tools can help you keep track of the schema and provide valuable information for developers and administrators. Additionally, some database management systems provide built-in support for generating schema documentation or exporting metadata.

## **view, user, constraints etc:**

It seems like you're interested in learning about database concepts such as views, users, and constraints. Let's briefly explain each of these concepts:

### **1. Views:**

- A view in a database is a virtual table that doesn't store data itself but presents data from one or more underlying tables in a structured manner. Views can be used to:
  - Simplify complex queries by abstracting the underlying structure.
  - Control access to data by limiting what columns or rows users can see.
  - Provide a customized perspective on the data.
- Views are defined by SQL queries and can be queried like regular tables.

### **2. Users:**

- In the context of a database management system, a user is an account that can access and interact with the database. Users have distinct usernames and may have their own permissions and roles.
- User accounts are used to control who can perform operations (e.g., querying, updating) on the database and what data they can access.
- User management is crucial for database security and access control.

### **3. Constraints:**

- Constraints are rules or conditions that are enforced on data in a database to maintain data integrity and consistency. Common types of constraints include:
  - Primary Key: Ensures that each row in a table has a unique identifier.
  - Foreign Key: Enforces referential integrity by linking one table's column to another table's primary key.
  - Unique Constraint: Ensures that values in a column (or combination of columns) are unique across rows.
  - Check Constraint: Verifies that values in a column meet specific conditions.
  - Not Null Constraint: Ensures that a column cannot contain null (empty) values.
- Constraints help maintain data quality and prevent incorrect or inconsistent data from being inserted or updated in the database.

## **Creating view:**

To create a view in SQL, you use the `CREATE VIEW` statement followed by the view name and the SQL query that defines the view. Here's the general syntax for creating a view:

```
CREATE VIEW view_name AS
SELECT column1, column2, ...
FROM table_name
WHERE conditions;
```

Here's a step-by-step guide to creating a view:

### **1. Determine the Data You Want in the View:**

Decide what data you want to include in your view. You will need to write a SQL query that retrieves this data from one or more tables.

### **2. Write the SQL Query for the View:**

Create the SQL query that retrieves the data for your view. This query can include conditions, joins, and any other SQL features you need to shape the data as desired.

### **3. Use the `CREATE VIEW` Statement:**

Use the `CREATE VIEW` statement to create the view. Provide a name for the view and specify the SQL query that defines the view. The view name should follow the naming conventions of your database system.

```
CREATE VIEW my_view AS
SELECT column1, column2
FROM my_table
WHERE condition;
```

In this example, `my\_view` is the name of the view, and it retrieves data from `my\_table` based on the specified conditions.

#### 4. Execute the SQL Statement:

Execute the `CREATE VIEW` statement in your SQL client or database management tool. If the statement is executed successfully, the view will be created.

#### 5. Query the View:

Once the view is created, you can query it just like you would query a regular table. Use the `SELECT` statement to retrieve data from the view:

```
SELECT FROM my_view;
```

This query will return the data as defined by the view's SQL query.

That's it! You've successfully created a view in your database. Keep in mind that views are virtual tables and do not store data themselves; they provide a convenient way to access and work with data from one or more underlying tables. Any changes made to the underlying tables will be reflected in the view's results.

### **Renaming the Column of a view, Updating, Selection, destroying view,**

In SQL, you can rename the columns of a view, update the view's definition, select data from a view, and destroy (drop) a view. Here's how you can perform each of these actions:

#### 1. Renaming the Column of a View:

To rename a column in a view, you need to recreate the view with the desired column names using the `AS` keyword in the view's SQL query. Here's an example of how to rename a column in a view:

```
CREATE OR REPLACE VIEW my_view AS
SELECT old_column_name AS new_column_name
FROM my_table;
```

In this example, `old\_column\_name` is renamed as `new\_column\_name` in the view.

#### 2. Updating the View's Definition:

To update the definition of an existing view, you can use the `CREATE OR REPLACE VIEW` statement to redefine the view. This is useful when you want to modify the query used by the view.

```
CREATE OR REPLACE VIEW my_view AS
SELECT column1, column2
FROM my_table
WHERE condition;
```

This statement replaces the existing definition of `my\_view` with the new SQL query.

### 3. Selecting Data from a View:

To select data from a view, you can use the `SELECT` statement as you would with a regular table. Simply specify the columns you want to retrieve in your query.

```
SELECT FROM my_view;
```

This query retrieves data from the `my\_view` view.

### 4. Dropping (Destroying) a View:

To delete (drop) a view from the database, you can use the `DROP VIEW` statement.

```
DROP VIEW my_view;
```

This statement removes the `my\_view` view from the database.

Keep in mind the following:

- When you recreate a view (as in renaming or updating the view), you need to ensure that the new query provides the same structure and data types for the columns as the original view, or you may need to update dependent queries or applications accordingly.
- Be cautious when dropping a view, as it cannot be undone. Ensure that you no longer need the view and that it won't impact other parts of your database or applications.
- The ability to rename, update, and drop views may vary depending on the database management system you are using, so consult your database system's documentation for specific syntax and capabilities.

## Indexes:

An index in a database management system (DBMS) is a data structure that improves the performance of database queries. Indexes are created on one or more columns in a table, and they store a sorted list of the values in those columns. When a query is executed, the DBMS can use the index to quickly find the rows that match the query criteria.

There are two main types of indexes in DBMSs:

- Clustered indexes:** A clustered index stores the data in the table in the same order as the index. This can improve performance for queries that retrieve large amounts of data in a specific order.
- Non-clustered indexes:** A non-clustered index stores the index values separately from the data in the table. This can improve performance for queries that retrieve a small number of rows from a large table.

## Creating Indexes,

## Syntax for Creating an Index:

```
CREATE INDEX index_name  
ON table_name (column1, column2, ...);
```

- ``index_name``: A user-defined name for the index.
- ``table_name``: The name of the table on which you want to create the index.
- ``(column1, column2, ...)``: The column(s) on which you want to create the index. You can create an index on one or more columns.

### Example of Creating a Simple Index:

Suppose you have a table called ``employees`` with a column ``employee_id`` that you frequently use in your queries. You can create an index on the ``employee_id`` column as follows:

```
CREATE INDEX idx_employee_id  
ON employees (employee_id);
```

## DCL: Granting Permissions:

In SQL, Data Control Language (DCL) statements are used to control access to database objects by granting or revoking permissions to users and roles. One of the key DCL operations is granting permissions to allow users to perform certain actions on database objects. Here's how you can grant permissions:

### Syntax for Granting Permissions:

```
GRANT permission(s) ON object TO user_or_role [WITH GRANT OPTION];
```

- ``permission(s)``: The specific permissions you want to grant, such as SELECT, INSERT, UPDATE, DELETE, EXECUTE, etc.
- ``object``: The database object (e.g., table, view, procedure) on which you want to grant the permissions.
- ``user_or_role``: The user or role to whom you want to grant the permissions.
- ``WITH GRANT OPTION`` (optional): Allows the user or role to grant the same permissions to others.

### Examples of Granting Permissions:

1. Granting SELECT permission on a table to a user:

```
GRANT SELECT ON employees TO user1;
```

This statement grants user1 permission to select (read) data from the ``employees`` table.

2. Granting INSERT and UPDATE permissions on a table to a role with the ability to grant those permissions to others:

```
GRANT INSERT, UPDATE ON products TO sales_role WITH GRANT OPTION;
```

This statement grants the sales\_role the ability to insert and update data in the ``products`` table and



allows the role to grant the same permissions to other users or roles.

3. Granting EXECUTE permission on a stored procedure to a user:

**GRANT EXECUTE ON my\_procedure TO user2;**

This statement allows user2 to execute the `my\_procedure` stored procedure.

### **Revoke permission - Creating and managing User:**

Revoking permissions in SQL is an essential part of managing user access to database objects. To revoke permissions previously granted to users or roles, you can use the `REVOKE` statement. Here's how to revoke permissions, as well as how to create and manage users:

1. Revoking Permissions:

Syntax for Revoking Permissions:

**REVOKE permission(s) ON object FROM user\_or\_role;**

- `permission(s)`: The specific permissions you want to revoke, such as SELECT, INSERT, UPDATE, DELETE, EXECUTE, etc.
- `object`: The database object (e.g., table, view, procedure) on which you want to revoke the permissions.
- `user\_or\_role`: The user or role from whom you want to revoke the permissions.

Examples of Revoking Permissions:

1. Revoking SELECT permission on a table from a user:

**REVOKE SELECT ON employees FROM user1;**

This statement revokes user1's permission to select (read) data from the `employees` table.

2. Revoking INSERT and UPDATE permissions on a table from a role:

**REVOKE INSERT, UPDATE ON products FROM sales\_role;**

This statement revokes the ability of the `sales\_role` to insert and update data in the `products` table.

3. Revoking EXECUTE permission on a stored procedure from a user:

**REVOKE EXECUTE ON my\_procedure FROM user2;**

This statement removes user2's ability to execute the `my\_procedure` stored procedure.

2. Creating and Managing Users:

User management is typically handled using Data Control Language (DCL) statements like `CREATE USER` and `ALTER USER`. The exact syntax and user management capabilities may vary depending on your specific database system. Here's a general guideline for creating and managing users:

Creating a User:

**CREATE USER username IDENTIFIED BY 'password';**

- 'username': The name of the user you want to create.
- 'password': The password for the user.

Example:

**CREATE USER user1 IDENTIFIED BY 'mypassword';**

Altering a User (Changing Password):

**ALTER USER username IDENTIFIED BY 'new\_password';**

- 'username': The name of the user whose password you want to change.
- 'new\_password': The new password for the user.

Example:

**ALTER USER user1 IDENTIFIED BY 'newpassword';**

Dropping a User:

**DROP USER username;**

- 'username': The name of the user you want to drop (remove) from the database.

Example:

**DROP USER user1;**

**Create, Role,  
provide privileges,  
dictionaries,**

It seems you want to create a role, provide privileges to that role, and understand the concept of dictionaries in the context of database management. Let's break down these topics:

1. Creating a Role:

In SQL, a role is a named collection of privileges that can be assigned to one or more users. Roles are used to simplify the management of permissions and access control. Here's how you can create a role:

**CREATE ROLE role\_name;**

- 'role\_name': The name you want to give to the role.

Example:

**CREATE ROLE sales\_team;**

This creates a role named 'sales\_team'.

## 2. Providing Privileges to a Role:

Once you have created a role, you can grant specific privileges to that role. Privileges can include permissions to access tables, execute procedures, and perform other database actions. Here's how to grant privileges to a role:

**GRANT privilege(s) ON object TO role\_name;**

- `privilege(s)`: The specific privileges you want to grant (e.g., SELECT, INSERT, EXECUTE).
- `object`: The database object (e.g., table, view, procedure) to which you want to grant privileges.
- `role\_name`: The name of the role to which you want to grant the privileges.

Example:

**GRANT SELECT, INSERT ON sales\_data TO sales\_team;**

This grants the `SELECT` and `INSERT` privileges on the `sales\_data` table to the `sales\_team` role.

## 3. Dictionaries in Database Management:

In the context of database management, dictionaries are typically associated with data dictionaries or metadata repositories. A data dictionary is a central repository of information about the data and database schema. It provides descriptions of tables, columns, data types, constraints, and other metadata.

Some databases have system tables or views that serve as data dictionaries. For example, in many database systems, the `information\_schema` database contains tables and views that store metadata about the database.

Here are some common uses of a data dictionary:

- Storing information about database objects (e.g., tables, views, columns).
- Documenting data types and constraints.
- Tracking dependencies between objects.
- Supporting data validation and integrity checks.

### Triggers:

A trigger is a set of instructions or a predefined action that automatically responds to specific events or changes in the database. Triggers are used to enforce data integrity, implement business rules, log changes, or perform other actions when certain database events occur. Here's a more detailed explanation of triggers:

#### 1. Event-Based Actions:

- Triggers are event-driven and execute automatically in response to specific events or actions in the database. Common events that trigger the execution of a trigger include INSERT, UPDATE, DELETE, and other data modification operations.

#### 2. Actions and Logic:

- A trigger consists of one or more SQL statements or a procedural code block (e.g., PL/SQL, T-SQL) that defines the action to be taken when the trigger event occurs.

- Triggers can perform a wide range of actions, such as validating data, enforcing referential integrity, logging changes, sending notifications, or executing more complex business logic.

### 3. Types of Triggers:

- There are two main types of triggers based on when they execute:
  - Before Triggers (or "BEFORE" triggers): These triggers execute before the triggering event. They are often used for data validation and modification.
  - After Triggers (or "AFTER" triggers): These triggers execute after the triggering event. They are often used for logging changes or performing post-event actions.

### 4. Benefits and Use Cases:

- Triggers are valuable for enforcing data consistency and maintaining data integrity. For example, they can prevent invalid data from being inserted, updated, or deleted.
- Triggers can automate tasks that would otherwise require manual intervention, making database management more efficient.
- They are often used to implement auditing, logging, and security features, such as tracking changes to critical data.

### 5. Challenges and Considerations:

- While triggers offer automation and data consistency benefits, they can also introduce complexity and potential performance overhead, especially if not well-designed or overused.
- Triggers should be used judiciously and documented carefully to ensure that their behavior is understood and maintained.

### 6. Examples:

- A "before insert" trigger can validate that certain fields are not empty before allowing a new record to be added to a table.
- An "after update" trigger can log changes to a history table whenever a specific column is modified.
- A "before delete" trigger can prevent the deletion of records related to other records in a parent-child relationship.

### 7. Database System Support:

- The syntax and capabilities of triggers may vary among different DBMSs. Popular database systems like Oracle, SQL Server, MySQL, PostgreSQL, and others support triggers, but the specifics can differ.

## **syntax,**

### **Enabling/Disabling Triggers**

Certainly! In SQL, triggers are special database objects that automatically respond to predefined events or actions within a database. Here's an explanation of trigger syntax and how to enable/disable them:

#### 1. Trigger Syntax:

The syntax for creating a trigger varies depending on the specific database system you're using. However, I'll provide a generic representation of the syntax to give you an idea of how triggers are typically defined:

```
CREATE [OR REPLACE] TRIGGER trigger_name  
{BEFORE | AFTER | INSTEAD OF} {INSERT | UPDATE | DELETE} ON table_name  
[FOR EACH ROW]  
BEGIN  
    -- Trigger logic here  
END;
```

- 'CREATE [OR REPLACE] TRIGGER': This statement is used to create a new trigger with the given

name. The 'OR REPLACE' option allows you to modify an existing trigger with the same name.

- 'trigger\_name': The user-defined name of the trigger.
- '{BEFORE | AFTER | INSTEAD OF}': Specifies when the trigger should execute. "BEFORE" triggers run before the triggering event, "AFTER" triggers run after, and "INSTEAD OF" triggers are used in views and handle certain types of operations.
- '{INSERT | UPDATE | DELETE}': Specifies the database action that triggers the execution of the trigger.
- 'ON table\_name': Indicates the table on which the trigger is defined.
- '[FOR EACH ROW]': Optional clause that specifies that the trigger operates on each row affected by the triggering event. This is used in row-level triggers.

Inside the trigger, you can include SQL statements or procedural code (e.g., PL/SQL, T-SQL) to perform actions in response to the event.

## 2. Enabling/Disabling Triggers:

In some database systems, you can enable or disable triggers on a table. This is useful when you want to temporarily suspend the action of a trigger without dropping it. The specific syntax for enabling or disabling triggers varies by database system, but here's a general representation:

- Enabling a Trigger:

```
ALTER TABLE table_name ENABLE TRIGGER trigger_name;
```

- Disabling a Trigger:

```
ALTER TABLE table_name DISABLE TRIGGER trigger_name;
```

- Enabling/Disabling All Triggers on a Table:

Some database systems allow you to enable or disable all triggers on a table at once. The syntax for this operation may differ from system to system.

For example, in SQL Server, you can enable or disable a trigger as follows:

- To enable:

```
ENABLE TRIGGER trigger_name ON table_name;
```

- To disable:

```
DISABLE TRIGGER trigger_name ON table_name;
```